

When Overlapping Windows Invent Predictability: A Negative Control for Evaluation Protocol in Time-Series Machine Learning

Joshua de Freitas

August 2026

Abstract

We train a short-horizon limit-order-book classifier on synthetic data in which future price movement is independent of every feature by construction. Under the standard evaluation recipe — sliding windows at stride one, a random train/validation split, and normalisation fitted on the full dataset — the model scores +26.8 points above its majority-class baseline at a 20-step horizon. Under a purged and embargoed split, on identical data with the same architecture and the same seeds, it scores −8.4 points — *below* baseline, which is the correct result when nothing is learnable. The 35-point difference is produced entirely by how the validation set was drawn. Both figures are means over five independent generator paths with twenty split seeds each (900 runs); across paths the leaky excess ranges +26.0 to +28.2 and the purged excess −14.6 to −3.3, so the sign of each is a property of the protocol rather than of the realisation.

We then derive a closed-form expression for the accuracy a given split geometry makes available. It depends only on the window stride, the label horizon, the labelling threshold and the train fraction; no property of the model enters it. It is an approximation rather than a proved bound — an oracle memoriser sits just below it at $h \in \{5, 10\}$ and just above it at $h = 20$ — and it tracks that oracle to within 0.56 points at every horizon, averaged over six independent generator paths. The derivation is further checked against 105 measurements spanning 20 independent generator paths, five window sizes, five thresholds and five train fractions.

The derivation is prescriptive. With stride s and horizon h the governing correlation is $\rho(s) = \max(0, (h - s)/h)$, which is exactly zero at $s \geq h$. Overlap leakage therefore does not merely diminish as stride increases — it disappears once stride reaches the horizon. We confirm this, and quantify its cost: the usable sample shrinks by a factor of s .

Separately, we find that the pipeline under study is not reproducible. Neither the split draw, the weight initialisation nor the batch order is seeded, so repeated runs of identical code on identical data differ materially. A model comparison decided by a margin smaller than that spread is indistinguishable from running the same model twice.

All figures regenerate from committed scripts.

1 Introduction

Data leakage — the presence of information in training that would not be available at prediction time — is a well documented cause of overstated machine learning results. Kapoor and Narayanan [2] surveyed seventeen fields, identified 329 affected papers, and produced an eight-type taxonomy, and tooling exists to detect some leakage classes automatically [5]. In financial time series specifically, purging and embargoing are established remedies [3], and the broader problem of backtest overfitting has been treated at length [1].

None of that is in dispute here, and none of it is what this paper contributes.

The difficulty with demonstrating leakage on real data is that manufactured skill and genuine skill are indistinguishable in the output. A practitioner shown that purged cross-validation lowers

their score has a ready and not unreasonable reply: purging is conservative, it discards legitimate short-horizon signal along with the leak, and the lower number reflects the remedy rather than an honest measurement. That objection cannot be settled on data where signal might exist.

It can be settled on data where signal cannot. We therefore construct a *negative control*: a synthetic order book whose mid price is an i.i.d. multiplicative random walk and whose order sizes are independent draws, with no mechanism connecting either to future price. Predictive skill is not merely absent — it is impossible. Every point of measured accuracy above the majority-class baseline is, by construction, an artifact of the evaluation procedure, and no appeal to destroyed signal is available.

This yields two results. The first is a measurement: how much predictability the standard recipe invents from nothing. The second, and the one we consider the contribution, is that this quantity has a closed form. Given the geometry of the windows and the correlation structure of the labels, the accuracy available to be manufactured can be computed before any model is trained — and the same expression states the condition under which it becomes zero.

2 The negative control

2.1 Generator

Mid prices follow

$$p_t = p_{t-1} (1 + \varepsilon_t), \quad \varepsilon_t \sim \mathcal{N}(0, \sigma^2) \text{ i.i.d.}, \quad \sigma = 3 \times 10^{-4}. \quad (1)$$

Three bid and three ask levels are placed around the mid at a spread of one to three ticks, and every level's size is an independent uniform integer draw. No size, spread or level price depends on any future price. Forward returns are independent of all features.

That this holds is enforced rather than asserted: a permutation test over non-overlapping spans, Bonferroni corrected across features, runs in continuous integration and fails if any feature acquires a detectable relationship with the forward return. The test has been observed to fail when a signal is deliberately introduced.

2.2 Windows and labels

Features are built into sliding windows of $W = 100$ rows at stride one: window t spans rows $t, \dots, t + W - 1$. The label is the sign of the forward return over horizon h , thresholded at $\theta = 5 \times 10^{-4}$:

$$y_t = \begin{cases} +1 & r_t > \theta \\ 0 & |r_t| \leq \theta \\ -1 & r_t < -\theta \end{cases} \quad r_t = \frac{p_{t+W-1+h} - p_{t+W-1}}{p_{t+W-1}}. \quad (2)$$

Two consequences of stride-one construction matter throughout. Adjacent windows share $W - 1$ of W input rows. And adjacent labels are signs of forward returns computed over spans that share $h - 1$ of h increments.

3 Protocols

Three splitting protocols are compared, holding architecture, data, optimiser and epoch count fixed.

Purging removes training windows whose raw-row span reaches into the validation region; the embargo additionally drops a buffer of validation windows after the cut [3].

Every run seeds the split draw, the weight initialisation and the batch order, so the study is bit-reproducible even though the pipeline it studies is not (Section 8).

Protocol	Construction	Overlap
<code>random_split</code>	shuffle window indices, then cut	1.00
<code>chronological</code>	train on earlier windows, validate on later	0.11
<code>purged_embargoed</code>	chronological, purged, plus embargo	0.00

Table 1: Overlap is the fraction of validation windows sharing at least one raw row with some training window.

Configuration span. Five independent generator paths $\times$ three protocols $\times$ three horizons $\times$ twenty split seeds: 900 runs. Every number reported below is a mean over generator paths, quoted with its across-path standard deviation, and every claim is additionally checked on each path separately (Section 4.1). The study is *not* varied across window size, labelling threshold, train fraction, optimiser, epoch count, row count or architecture; Section 7 varies the first three for the derivation alone, and Section 10 states what that leaves open.

Accuracy is reported as **excess over the majority-class baseline computed on the same validation set**. The baseline moves from 0.575 to 0.398 across horizons because the class balance shifts with h , so raw accuracies are not comparable between rows.

4 Results

Protocol	h	Accuracy	Baseline	Excess	sd (paths)	range over paths
<code>random_split</code>	5	0.5383	0.5445	−0.0061	0.0071	[−0.0127, +0.0054]
<code>random_split</code>	10	0.5542	0.3953	+ 0.1589	0.0184	[+0.1338, +0.1828]
<code>random_split</code>	20	0.6446	0.3762	+ 0.2685	0.0080	[+0.2603, +0.2818]
<code>chronological</code>	5	0.4862	0.5518	−0.0656	0.0144	[−0.0827, −0.0494]
<code>chronological</code>	10	0.3412	0.4153	−0.0741	0.0254	[−0.0972, −0.0342]
<code>chronological</code>	20	0.3295	0.4123	−0.0828	0.0343	[−0.1295, −0.0474]
<code>purged_embargoed</code>	5	0.4897	0.5504	−0.0606	0.0202	[−0.0817, −0.0357]
<code>purged_embargoed</code>	10	0.3324	0.4167	−0.0843	0.0274	[−0.1035, −0.0380]
<code>purged_embargoed</code>	20	0.3280	0.4125	−0.0845	0.0449	[−0.1464, −0.0327]

Table 2: Validation-only accuracy, averaged over five independent generator paths with twenty split seeds each. *Excess* is accuracy minus the majority-class baseline on the same validation set; *sd (paths)* is the standard deviation of that excess across paths and the final column its full range. Only `random_split` at $h \geq 10$ clears its baseline, and it clears it on every path.

Three observations.

Only the leaky protocol shows skill. Every honest cell falls below its own baseline. A model trained on noise underperforms a constant predictor, which is the correct outcome.

Overlap fraction is not the predictor. `chronological` sits at 0.11 overlap and behaves like `purged_embargoed` at 0.00, not like something one-ninth of the way to `random_split`. A chronological cut leaves a thin seam and manufactures nothing. What matters is whether near-duplicate *neighbours* are separated across the split, which shuffling guarantees and a time cut prevents.

Leakage widens as well as inflates. Within a single generator path, the leaky $h = 20$ cell disperses far more across split seeds than its purged counterpart: on the reference path,




Figure 1: Two cells of nine clear their own baseline, and both have overlap 1.00.

standard deviation 0.0775 against 0.0421, with individual seeds spanning 0.468 to 0.751. Note the contrast with the across-path standard deviation of 0.0080 in the table above — the manufactured advantage is an order of magnitude more stable across data realisations than across initialisations, which is what one expects of a property of the evaluation geometry rather than of any fitted model.

4.1 Every generator path, separately

Table 2 averages over paths, which can hide a conclusion that holds only on average. It does not here. The same nine cells, unpooled:

Protocol	h	Path 0	Path 1	Path 2	Path 3	Path 4
random_split	5	-0.0070	-0.0049	-0.0113	-0.0127	+0.0054
random_split	10	+0.1497	+0.1663	+0.1338	+0.1617	+0.1828
random_split	20	+0.2603	+0.2682	+0.2671	+0.2651	+0.2818
chronological	5	-0.0494	-0.0522	-0.0686	-0.0752	-0.0827
chronological	10	-0.0649	-0.0342	-0.0829	-0.0972	-0.0913
chronological	20	-0.1076	-0.0669	-0.0623	-0.0474	-0.1295
purged_embargoed	5	-0.0457	-0.0357	-0.0817	-0.0791	-0.0609
purged_embargoed	10	-0.0808	-0.0380	-0.0975	-0.1018	-0.1035
purged_embargoed	20	-0.1027	-0.0904	-0.0503	-0.0327	-0.1464

Table 3: Excess over the majority baseline, by generator path. Every honest cell is negative on every path; **random_split** is positive on every path at $h \geq 10$. No path contradicts the conclusion.

Nine cells $\times$ five paths is 45 signs, and all 45 agree with Table 2. The conclusion does not rest on a choice of generator seed.

5 The ceiling

The preceding section measures what one architecture achieves. A different question admits a closed-form answer: given the geometry alone, how much accuracy is *available* to be manufac-

tured?

Adjacent windows have forward returns that are sums of h i.i.d. increments sharing $h - 1$ of them, so

$$\rho = \text{corr}(r_t, r_{t+1}) = \frac{h-1}{h}, \quad (3)$$

which increases with the horizon. Writing $z = \theta/(\sigma\sqrt{h})$ for the threshold in units of the h -step return’s standard deviation, and Φ_2 for the standard bivariate normal CDF at correlation ρ , the probability that two adjacent windows carry the same label is

$$A(h) = \underbrace{\Phi_2(-z, -z)}_{\text{both down}} + \underbrace{\Phi_2(z, z) - 2\Phi_2(-z, z) + \Phi_2(-z, -z)}_{\text{both flat}} + \underbrace{1 - 2\Phi(z) + \Phi_2(z, z)}_{\text{both up}}. \quad (4)$$

A memoriser copies the label of its *nearest* training window, which is not always the adjacent one. Under a random split at train fraction f , the distance d to the nearest training neighbour satisfies

$$\Pr(d > k) = (1 - f)^{2k}, \quad \Pr(d = k) = (1 - f)^{2(k-1)} - (1 - f)^{2k}, \quad (5)$$

and a neighbour d windows away shares $h - d$ of h increments, so $\rho_d = \max(0, (h - d)/h)$. The ceiling is the expectation of the agreement probability over that distance distribution:

$$C(h) = \sum_{d \geq 1} \Pr(d) A(\rho_d). \quad (6)$$

Nothing about the model appears in (6). It is computable before any training run.

An earlier version of this derivation collapsed every case beyond $d = 1$ to the majority class. That is wrong, and diagnosably so: it under-predicted by 0.094 at $h = 20$, $f = 0.5$ — where a quarter of validation windows have no adjacent twin — with the error decaying monotonically to 0.004 at $f = 0.9$. The train-fraction sweep in Section 7 was what exposed it.

5.1 Verification

We evaluate an *oracle*: for each validation window, predict the label of the temporally nearest training window. The neighbour is supplied rather than learned, which isolates the correlation model from the separate question of whether a learner could locate its twin.

h	ρ	Ceiling $C(h)$	Oracle (6 paths)	Error
5	0.800	0.6896	0.6873 ± 0.0044	-0.0023
10	0.900	0.7479	0.7465 ± 0.0073	-0.0014
20	0.950	0.8093	0.8149 ± 0.0085	+0.0056

Table 4: The oracle column is the mean over six independent generator paths, of the mean over twenty split seeds within each, $\pm$ the standard deviation of those six path means. The prediction error is smaller than that spread at every horizon, and no parameter is fitted. Note the sign: the closed form slightly over-predicts at $h \in \{5, 10\}$ and under-predicts at $h = 20$, consistent with the class-balance term $M(h)$ discussed in Section 10. It is an approximation to a bound, not a bound that has been proved to hold pathwise.

The expression binds only under a random split. Under **chronological** and **purged_embargoed** the nearest training window is the *same* window for every validation example, so the strategy degenerates to a constant predictor and scores at or below baseline: across 36 measurements — six generator paths $\times$ three horizons $\times$ two honest protocols — the oracle’s excess over baseline runs from -0.38 to exactly 0.00 and is positive in none of them. Copying one’s neighbour is useful only when the neighbour has been scattered into the training set.

5.2 Availability is not exploitation




Figure 2: The shaded region is the leak the split makes available. The trained network takes none of it at $h = 5$ and 62% at $h = 20$.

Measured accuracy can be read as the fraction of available leak a given model recovers. Averaged over generator paths: none at $h = 5$, 45% at $h = 10$, 62% at $h = 20$. At $h = 5$ there are 14.5 points of leak available and this architecture takes none of them — it is 0.6 points *below* its baseline, which is within the across-path spread of 0.7 points.

The ceiling therefore describes the *exposure* a split geometry creates. Whether a particular model converts that exposure into a headline number is a separate, model-dependent question. Both are worth knowing, and only the first can be computed in advance.

For completeness: a 1-nearest-neighbour classifier in raw feature space scores 0.29–0.45 across 360 runs on six generator paths, below baseline everywhere. Euclidean distance over price-level features locates windows at similar *price levels*, not temporal twins. The trained network’s performance is therefore not naive distance matching.

6 The remedy

Equation (3) assumed stride one. With stride s , retained windows are s rows apart and their forward returns share $h - s$ of h increments:

$$\rho(s) = \max\left(0, \frac{h - s}{h}\right). \quad (7)$$

At $s \geq h$ this is exactly zero: retained windows share no increments, their labels are independent, and copying a neighbour cannot beat the baseline. Overlap leakage does not diminish asymptotically — it terminates.

Two consequences deserve separation.

The rule. Set stride $\geq$ horizon and overlap leakage is gone rather than reduced. At intermediate strides (7) and (6) quantify the remaining exposure, so the choice can be made numerically.

Leakage is predictable, and it vanishes once stride reaches the horizon




Figure 3: Closed form (line) against measured oracle (points) across the stride sweep. The dotted vertical marks $s = h$.

h	Stride	$\rho(s)$	Ceiling	Oracle	Excess over majority	n windows
20	1	0.950	0.8093	0.8129	+0.4449	19881
20	2	0.900	0.7359	0.7412	+0.3721	9941
20	5	0.750	0.6114	0.6139	+0.2467	3977
20	10	0.500	0.4903	0.4971	+0.1287	1989
20	15	0.250	0.4067	0.4018	+0.0353	1326
20	20	0.000	0.3361	0.3380	-0.0337	995

Table 5: Mean over five independent generator paths, 20,000 rows. The excess over baseline is positive at every stride below the horizon and negative at $s = h$, on every path.

The cost. Stride s divides the usable sample by s : 4881 windows become 245 at $h = 20$. This is the actual trade — leakage against sample size — and it is why stride-one windowing is common. The formula does not remove the trade; it prices it.

Beyond $s = h$ the memoriser scores *below* baseline, because with independent labels copying a neighbour yields $\sum_i p_i^2 < \max_i p_i$. This is predicted behaviour and serves as an independent check on the derivation.

The verification is weakest where the remedy is strongest. Raising the stride divides the sample by s , which reduces the validation set the remedy is verified on by the same factor. At 5,000 rows and $h = s = 20$ only 49 validation windows remain, so a single window is worth 2.04 accuracy points; on one of five generator paths the measured excess came out at +0.043 — two windows — against a claim that it should be non-positive. At 20,000 rows the same cell has 199 validation windows and all five paths are negative. Pooling all 255 measured cells, the closed form’s error against the oracle falls monotonically with validation size: 0.024 mean below 60 windows, 0.0065 above 400. Restricted to cells with at least 150 validation windows the mean error is 0.0095 and the maximum 0.051.

The apparent breakdown is therefore a property of the measurement, not of the remedy. It is reported because a single-path run at 5,000 rows would have shown a clean result and concealed it.

7 Robustness of the derivation

The ceiling contains stride, horizon, threshold, per-step volatility and train fraction. It does *not* contain the window size W , which is a falsifiable prediction: widening the feature window changes what the model sees but not how correlated the labels of neighbouring windows are, and the leak lives in the labels.

Sweep	Range	Max error	Mean error
Generator path	20 independent seeds	0.0208	0.0073
Window size	$W \in \{25, \dots, 400\}$	0.0053	0.0031
Threshold	$\theta \in [10^{-4}, 2 \times 10^{-3}]$	0.0223	0.0081
Train fraction	$f \in [0.5, 0.9]$	0.0051	0.0022

Table 6: Closed form against measured oracle, 105 configurations.

A sixteen-fold change in window size moves the oracle by less than 0.006 while the prediction is held fixed. The prediction survives. Twenty independent generator paths give a mean absolute error of 0.007, so the result is not a property of one realisation.

8 Reproducibility of the pipeline under study

The pipeline seeds neither the split draw, the weight initialisation, nor the batch order. Repeated runs of identical code on identical data therefore differ. Six unseeded repetitions of one configuration (`random_split`, $h = 20$, generator seed 42) yield validation accuracies spanning 0.5036 to 0.7134 (mean 0.6395, sd 0.0830) and all-window accuracies spanning 0.5206 to 0.7408 (mean 0.6712, sd 0.0864); runs and metrics are committed in `results/study-multiseed/summary.json` under `unseeded_arm`. An earlier estimate of 0.67–0.75 predated the committed harness and is superseded by this artifact.

Stated separately from the leakage result because it is independent and applies more broadly: **a model comparison decided by a margin smaller than the pipeline’s own run-to-run spread is indistinguishable from running the same model twice.** This includes the architecture comparison recorded in the source repository. Seed sensitivity of this magnitude is consistent with reports elsewhere [4].

Three lines of seeding render the pipeline bit-deterministic.

9 The normalisation leak, measured

The pipeline also fits its feature normalisation on the whole dataset: each size column is divided by its maximum over all rows, computed before any split exists. An earlier version of this paper described that leak as "presumed small, not measured." It is neither.

Replacing it with train-only z -scoring — fitted on $X[\text{train}]$ alone, then applied to both arms — and pairing on generator path, split seed and split, so that only the scaling differs:

The direction is the opposite of the one assumed. **Removing the normalisation leak makes the leaky protocol perform better**, by 4.3 points at $h=10$ and 7.3 at $h=20$. Dividing by a global maximum compresses the size features into a narrow band and flattens their structure; standardising each feature to unit variance preserves it, and a model that can see more structure identifies its near-duplicate twin more easily.

So the two defects were never independent contributions to be summed. One was partially masking the other, and the honest reading is:

Correct feature scaling amplifies a leak. It does not create one.

Cell	train-only – global-max (pp)	95% CI
random_split $h=20$	+7.26	[+4.90, +9.62]
random_split $h=10$	+4.30	[+2.44, +6.16]
random_split $h=5$	+0.48	[−0.18, +1.15]
chronological $h=5$	−1.91	[−3.51, −0.31]
chronological $h=10$	+0.06	[−1.15, +1.26]
chronological $h=20$	−0.44	[−1.88, +1.00]
purged_embargoed $h=5$	−2.72	[−3.97, −1.47]
purged_embargoed $h=10$	+0.95	[−0.53, +2.42]
purged_embargoed $h=20$	−0.00	[−1.61, +1.61]

Table 7: 50 matched pairs per cell across five generator paths.

Where a leak exists to exploit, better scaling exploits more of it. Where none exists — the honest protocols — the effect is null or slightly negative, and in no case does it lift a cell above baseline. **All 36 purged combinations, six generator paths $\times$ three horizons $\times$ two scalings, remain below their validation-majority baseline**; the narrowest margin is -2.69 points.

One consequence for the headline: with correct train-only scaling, **random_split** at $h=20$ reports 72.1% against a 39.8% baseline. Fixing one methodological error makes the other one worse.

And $h=5$ is not a negative control. An earlier draft treated the $h=5$ cell — where the leaky protocol sits at its baseline — as evidence that overlap alone is insufficient. It is not evidence of anything. Across ten path-and-scaling combinations, nine excesses are statistically indistinguishable from zero, the single exception is what ten comparisons produce by chance, and the sign varies by generator path. The correct statement is that at $h=5$ the effect is too small to resolve at this sample size, and we can neither confirm nor exclude it.

10 Limitations

- **No published result is shown to be incorrect.** We show what a procedure can manufacture on data without signal. Whether any particular result suffers the same effect requires that result’s data and code.
- **This is not a measurement of real markets.** The generator is a random walk without microstructure — appropriate for a negative control, a limitation for external validity.
- **The normalisation leak is now measured** rather than presumed, and it does not behave additively — see Section 6. The three defects are not separable channels and the paper does not decompose them as such.
- **The ceiling assumes i.i.d. increments.** Real returns exhibit volatility clustering and heavy tails; $\rho = (h - s)/h$ is exact only under the generator used here.
- **Predicted class balance** matches measurement within 0.3 standard deviations at $h \in \{5, 10\}$ across 20 independent paths, and under-predicts the majority share by approximately 1.3 standard deviations at $h = 20$. This affects the $M(h)$ term in (6), not the verified $A(h)$ term.
- **One generator family and one architecture.** A second model family is untested, so the fraction of the available leak a model recovers may be architecture-specific.

- **Configuration span.** The protocol comparison spans five generator paths $\times$ twenty split seeds; the ceiling six paths; the stride sweep five paths $\times$ two generator lengths; the normalisation comparison five paths $\times$ ten split seeds, plus a twenty-seed reference path. Not varied: window size in the trained arms, labelling threshold, train fraction, optimiser, epoch count, architecture.

11 Conclusion

On data constructed to contain no predictable structure, the standard evaluation recipe for time-series machine learning reports 64% accuracy against a 38% baseline. The honest protocol reports 33% against 41%. Averaged over five generator paths and twenty split seeds each, and with the same sign on every one of those paths, the difference is attributable entirely to the evaluation design; because the data contains no signal, the usual defence of the leaky protocol is unavailable.

The quantity has a closed form. The accuracy a split geometry makes available depends on stride, horizon, threshold and train fraction, and on nothing about the model. That expression is verified against measurement across 105 configurations spanning twenty generator paths, five window sizes, five thresholds and five train fractions, and it states its own remedy: at stride greater than or equal to the label horizon, the leak is zero.

Reproducing

```
python3 -m venv .venv && .venv/bin/pip install -e ".[test]"
.venv/bin/python run_study.py          # 900 seeded runs, 5 generator paths
.venv/bin/python run_ceiling.py        # closed form, oracle, 1-NN
.venv/bin/python run_stride.py         # stride sweep
.venv/bin/python make_figures.py       # figures from committed results
```

Per-run evidence is committed under `results/`. Environment: Python 3.12.13, torch 2.13.0, numpy 2.5.1, pandas 3.0.5.

References

- [1] Bailey, D. H., Borwein, J. M., López de Prado, M., Zhu, Q. J. (2014). Pseudo-Mathematics and Financial Charlatanism: The Effects of Backtest Overfitting on Out-of-Sample Performance. *Notices of the American Mathematical Society*, 61(5), 458–471.
- [2] Kapoor, S., Narayanan, A. (2023). Leakage and the Reproducibility Crisis in ML-based Science. *Patterns*, 4(9), 100804. arXiv:2207.07048.
- [3] López de Prado, M. (2018). *Advances in Financial Machine Learning*. John Wiley & Sons.
- [4] Picard, D. (2021). `torch.manual_seed(3407)` is all you need: On the influence of random seeds in deep learning architectures for computer vision. arXiv:2109.08203.
- [5] Talagala, T. S. (2024). `tsdataleaks`: An R Package to Detect Potential Data Leaks in Forecasting Competitions. arXiv:2402.10522.